

Reinforcement Learning

申思杰

Outline

- Reinforcement Learning Background (model-free methods)
 - Value-based: **SARSA, Q-learning, Deep Q Network**
 - Policy-based: **REINFORCE, Actor-Critic**
- Asynchronous RL Framework
 - Asynchronous SARSA, Asynchronous Q-learning
 - Asynchronous Advantage Actor-Critic (A3C)

Markov Decision Process

- Markov Decision Process:

Definition

A *Markov Decision Process* is a tuple $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$

- \mathcal{S} is a finite set of states
- \mathcal{A} is a finite set of actions
- \mathcal{P} is a state transition probability matrix,
 $\mathcal{P}_{ss'}^a = \mathbb{P}[S_{t+1} = s' \mid S_t = s, A_t = a]$
- \mathcal{R} is a reward function, $\mathcal{R}_s^a = \mathbb{E}[R_{t+1} \mid S_t = s, A_t = a]$
- γ is a discount factor $\gamma \in [0, 1]$.

Markov Decision Process

- **Policy:** A policy π is a distribution over actions given states

$$\pi(a | s) = \mathbb{P}[A_t = a | S_t = s]$$

- **Episode:**

- At state S_t
- Perform action A_t according to policy π
- Get reward R_{t+1} and get to state S_{t+1}
- Loop until terminate or until time-step T

$$S_1, A_1, R_2, S_2, A_2, R_3, S_3, \dots$$

Markov Decision Process

- Return: Total discounted reward from time step t

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

- Value function:

- State-value function: $V_{\pi}(s) = \mathbb{E}_{\pi}[G_t | S_t = s]$

- Action-value function: $Q_{\pi}(s, a) = \mathbb{E}[G_t | S_t = s, A_t = a]$

Markov Decision Process

- Value function can be decomposed into immediate reward plus discounted value of successor state

$$V_{\pi}(s) = \mathbb{E}_{\pi}[R_{t+1} + \gamma V_{\pi}(S_{t+1}) | S_t = s]$$

$$Q_{\pi}(s, a) = \mathbb{E}[R_{t+1} + \gamma Q_{\pi}(S_{t+1}, A_{t+1}) | S_t = s, A_t = a]$$

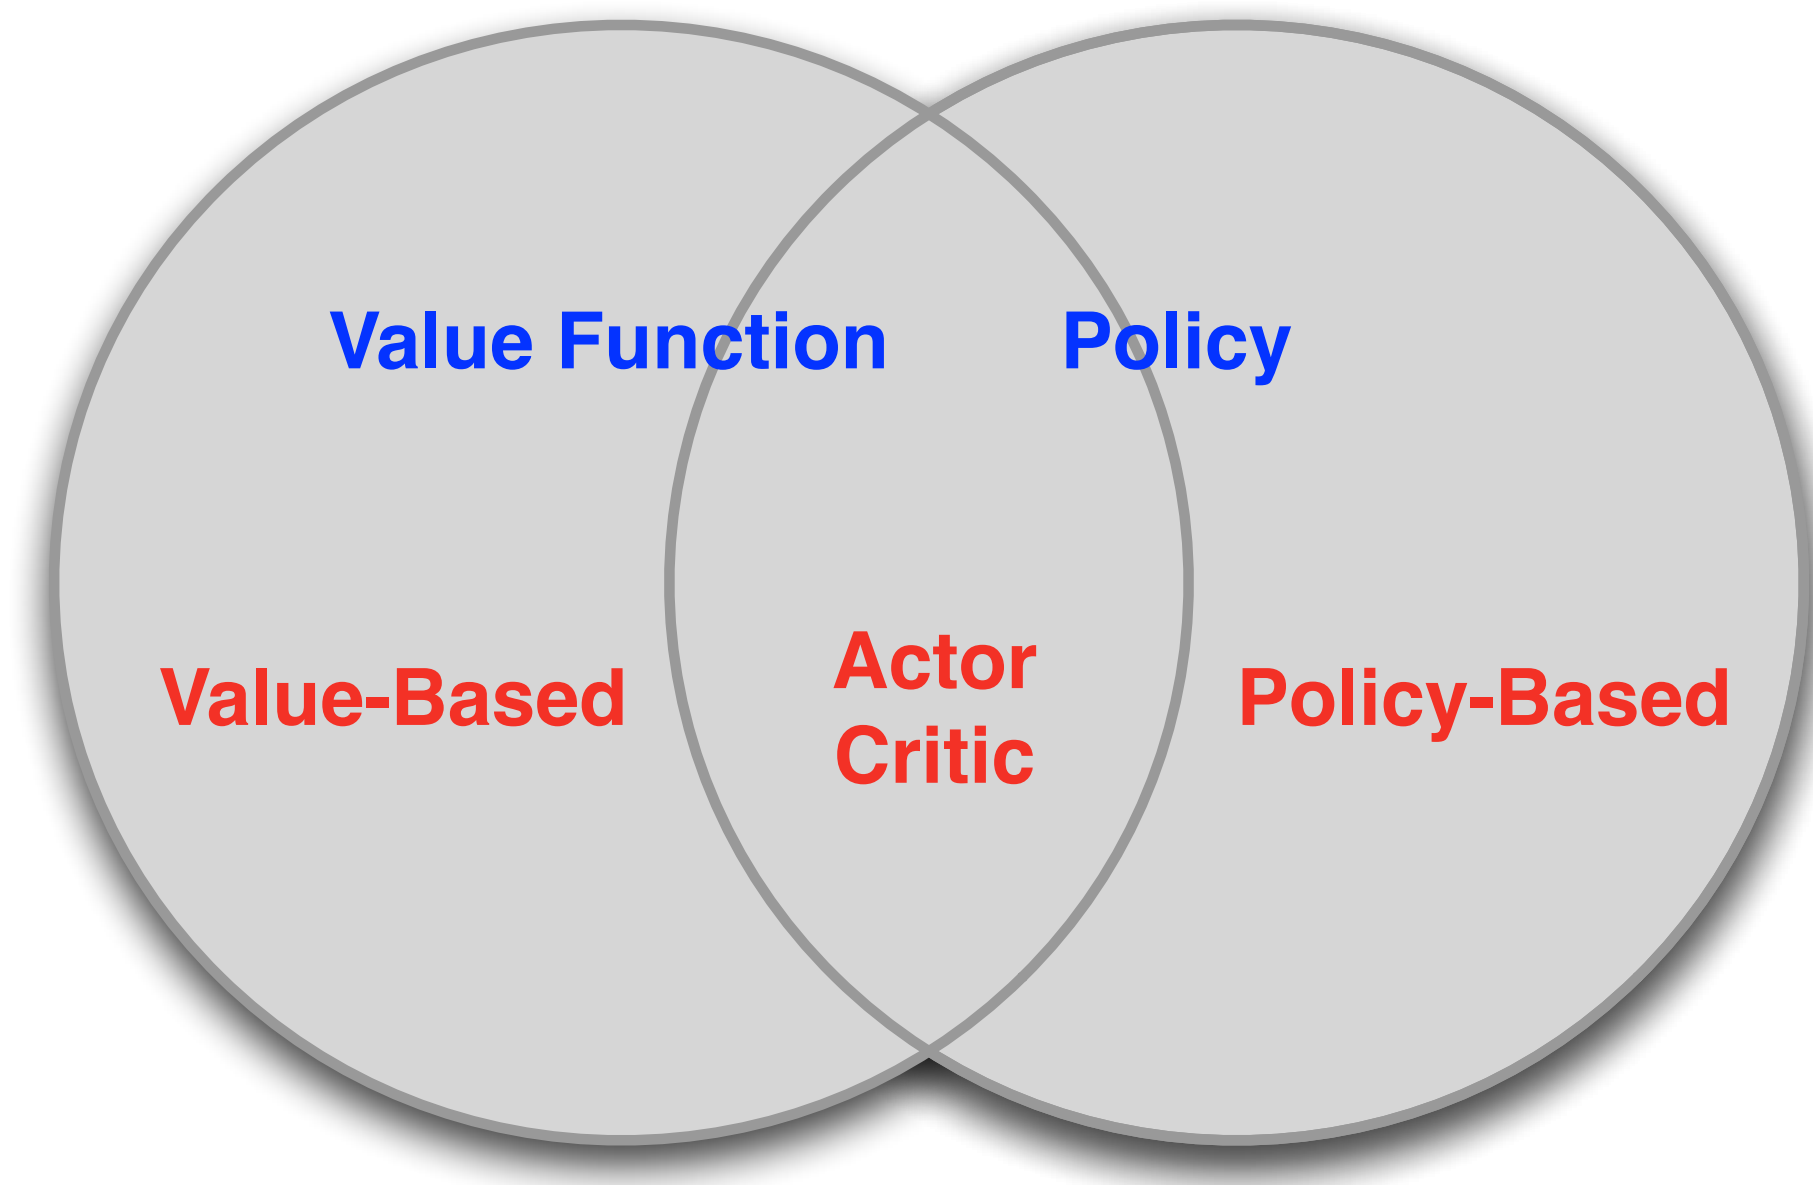
- Optimal value function:

$$V^*(s) = \max_{\pi} V_{\pi}(s)$$

$$Q^*(s, a) = \max_{\pi} Q_{\pi}(s, a)$$

An optimal policy can be found by maximizing over $Q^*(s, a)$

Model-Free Reinforcement Learning Methods



Value-based Methods

- Directly approximate the optimal action value function $Q^*(s, a)$
 - e.g. Using a table to store the approximation of $Q^*(s, a)$
 - e.g. Using neural network $Q(s, a; \theta)$ with parameters θ to fit $Q^*(s, a)$
- Policy can be derived from the learned $Q(s, a; \theta)$
 - e.g. ϵ -greedy policy

Value-based Methods

- Monte-Carlo control:

Policy evaluation Monte-Carlo policy evaluation, $Q \approx q_\pi$

Policy improvement ϵ -greedy policy improvement

- On-policy and off-policy learning:

- On-policy learning

- “Learn on the job”

- Learn about policy π from experience sampled from π

- Off-policy learning

- “Look over someone’s shoulder”

- Learn about policy π from experience sampled from μ

Value-based Methods

- SARSA for on-policy control:

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$
Repeat (for each episode):
 Initialize S
 Choose A from S using policy derived from Q (e.g., ϵ -greedy)
 Repeat (for each step of episode):
 Take action A , observe R, S'
 Choose A' from S' using policy derived from Q (e.g., ϵ -greedy)
 $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$
 $S \leftarrow S'; A \leftarrow A';$
 until S is terminal

Value-based Methods

- Q-learning for off-policy control


Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$
Repeat (for each episode):
 Initialize S
 Repeat (for each step of episode):
 Choose A from S using policy derived from Q (e.g., ϵ -greedy)
 Take action A , observe R, S'
 $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$
 $S \leftarrow S'$
 until S is terminal

Value-based Methods

- SARSA:

- Policy: ϵ -greedy


$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha (R + \gamma Q(s', a'))$$



- Q-learning:


- Target policy: greedy; Behavior policy: ϵ -greedy

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + \alpha \left(R + \gamma \max_{a'} Q(s', a') \right)$$




Value-based Methods

- Deep Q Network (DQN):



Tradition Methods



Deep Methods

Value-based Methods

DQN uses **experience replay** and **fixed Q-targets**

- Take action a_t according to ϵ -greedy policy
- Store transition $(s_t, a_t, r_{t+1}, s_{t+1})$ in replay memory \mathcal{D}
- Sample random mini-batch of transitions (s, a, r, s') from \mathcal{D}
- Compute Q-learning targets w.r.t. old, fixed parameters w^-
- Optimise MSE between Q-network and Q-learning targets

$$\mathcal{L}_i(w_i) = \mathbb{E}_{s,a,r,s' \sim \mathcal{D}_i} \left[\left(r + \gamma \max_{a'} Q(s', a'; w_i^-) - Q(s, a; w_i) \right)^2 \right]$$

- Using variant of stochastic gradient descent

Policy-based Methods

- Directly parameterize the policy $\pi_\theta(a | s)$, find the best θ
- Measure the quality of policy: policy objective functions

- Start value (episodic):

$$J_1(\theta) = V^{\pi_\theta}(s_1) = \mathbb{E}_{\pi_\theta}[v_1]$$

- Average value (continuing):

$$J_{avV}(\theta) = \sum_s d^{\pi_\theta}(s) V^{\pi_\theta}(s)$$

- Average reward per time-step (continuing):

$$J_{avR}(\theta) = \sum_s d^{\pi_\theta}(s) \sum_a \pi_\theta(s, a) \mathcal{R}_s^a$$

d^{π_θ} is stationary distribution of Markov chain for π_θ

Policy-based Methods

- Policy gradient to update parameters:

$$\Delta\theta = \alpha \nabla_{\theta} J(\theta)$$

$\nabla_{\theta} J(\theta)$ is called policy gradient

- Compute policy gradient: Likelihood ratios:

$$\begin{aligned}\nabla_{\theta} \pi_{\theta}(s, a) &= \pi_{\theta}(s, a) \frac{\nabla_{\theta} \pi_{\theta}(s, a)}{\pi_{\theta}(s, a)} \\ &= \pi_{\theta}(s, a) \nabla_{\theta} \log \pi_{\theta}(s, a)\end{aligned}$$

Policy-based Methods

- An example: Gradient of one-step MDPs:

$$\begin{aligned} J(\theta) &= \mathbb{E}_{\pi_{\theta}} [r] \\ &= \sum_{s \in \mathcal{S}} d(s) \sum_{a \in \mathcal{A}} \pi_{\theta}(s, a) \mathcal{R}_{s,a} \\ \nabla_{\theta} J(\theta) &= \sum_{s \in \mathcal{S}} d(s) \sum_{a \in \mathcal{A}} \pi_{\theta}(s, a) \nabla_{\theta} \log \pi_{\theta}(s, a) \mathcal{R}_{s,a} \\ &= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) r] \end{aligned}$$

Policy-based Methods

- Policy Gradient Theorem:

Theorem

*For any differentiable policy $\pi_\theta(s, a)$,
for any of the policy objective functions $J = J_1, J_{avR}$, or $\frac{1}{1-\gamma} J_{avV}$,
the policy gradient is*

$$\nabla_\theta J(\theta) = \mathbb{E}_{\pi_\theta} [\nabla_\theta \log \pi_\theta(s, a) Q^{\pi_\theta}(s, a)]$$

Policy-based Methods

- REINFORCE (Monte-Carlo Policy Gradient):
 - Using return v_t as a sample of $Q^{\pi_\theta}(s_t, a_t)$

```
function REINFORCE
  Initialise  $\theta$  arbitrarily
  for each episode  $\{s_1, a_1, r_2, \dots, s_{T-1}, a_{T-1}, r_T\} \sim \pi_\theta$  do
    for  $t = 1$  to  $T - 1$  do
       $\theta \leftarrow \theta + \alpha \nabla_\theta \log \pi_\theta(s_t, a_t) v_t$ 
    end for
  end for
  return  $\theta$ 
end function
```

Policy-based Methods

- Reduce variance using a critic:

- We use a **critic** to estimate the action-value function,

$$Q_w(s, a) \approx Q^{\pi_\theta}(s, a)$$

- Actor-critic algorithms maintain *two* sets of parameters

Critic Updates action-value function parameters w

Actor Updates policy parameters θ , in direction suggested by critic

- Actor-critic algorithms follow an *approximate* policy gradient

$$\nabla_\theta J(\theta) \approx \mathbb{E}_{\pi_\theta} [\nabla_\theta \log \pi_\theta(s, a) Q_w(s, a)]$$

$$\Delta\theta = \alpha \nabla_\theta \log \pi_\theta(s, a) Q_w(s, a)$$

Policy-based Methods

- Actor-Critic:

- Using linear value function to approximate $Q_w(s, a) = \phi(s, a)^T w$

```
function QAC
  Initialise  $s, \theta$ 
  Sample  $a \sim \pi_\theta$ 
  for each step do
    Sample reward  $r = \mathcal{R}_s^a$ ; sample transition  $s' \sim \mathcal{P}_s^a$ .
    Sample action  $a' \sim \pi_\theta(s', a')$ 
     $\delta = r + \gamma Q_w(s', a') - Q_w(s, a)$ 
     $\theta = \theta + \alpha \nabla_\theta \log \pi_\theta(s, a) Q_w(s, a)$ 
     $w \leftarrow w + \beta \delta \phi(s, a)$ 
     $a \leftarrow a', s \leftarrow s'$ 
  end for
end function
```

Policy-based Methods

- Reduce variance using a baseline:

$$\begin{aligned}\mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) B(s)] &= \sum_{s \in \mathcal{S}} d^{\pi_{\theta}}(s) \sum_a \nabla_{\theta} \pi_{\theta}(s, a) B(s) \\ &= \sum_{s \in \mathcal{S}} d^{\pi_{\theta}} B(s) \nabla_{\theta} \sum_{a \in \mathcal{A}} \pi_{\theta}(s, a) \\ &= 0\end{aligned}$$

- Advantage Actor-Critic: Using state-value function as a baseline:

- Advantage function: $A^{\pi_{\theta}}(s, a)$

$$\begin{aligned}A^{\pi_{\theta}}(s, a) &= Q^{\pi_{\theta}}(s, a) - V^{\pi_{\theta}}(s) \\ \nabla_{\theta} J(\theta) &= \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) A^{\pi_{\theta}}(s, a)]\end{aligned}$$

Outline

- Reinforcement Learning Background (model-free methods)
 - Value-based: **SARSA, Q-learning, Deep Q Network**
 - Policy-based: **REINFORCE, Actor-Critic**
- Asynchronous RL Framework
 - Asynchronous SARSA, Asynchronous Q-learning
 - Asynchronous Advantage Actor-Critic (A3C)

Asynchronous Methods for Deep Reinforcement Learning

Volodymyr Mnih¹

Adrià Puigdomènech Badia¹

Mehdi Mirza^{1,2}

Alex Graves¹

Tim Harley¹

Timothy P. Lillicrap¹

David Silver¹

Koray Kavukcuoglu¹

¹ Google DeepMind

² Montreal Institute for Learning Algorithms (MILA), University of Montreal

VMNIH@GOOGLE.COM

ADRIAP@GOOGLE.COM

MIRZAMOM@IRO.UMONTREAL.CA

GRAVESA@GOOGLE.COM

THARLEY@GOOGLE.COM

COUNTZERO@GOOGLE.COM

DAVIDSILVER@GOOGLE.COM

KORAYK@GOOGLE.COM